

FrozenLake

January 15, 2022

0.1 Mount the Google Drive onto the Colab as the storage location.

Following the instructions returned from the below cell. You will click a web link and select the google account you want to mount, then copy the authorization code to the blank, press enter.

```
[1]: # This must be run within a Google Colab environment
from google.colab import drive
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call `drive.mount("/content/gdrive", force_remount=True)`.

0.2 Append the directory location where you upload the start code folder (In this problem, *RLalgs*) to the `sys.path`

E.g. `dir = '/content/drive/My Drive/RL/'`, start code folder is inside “RL” folder.

```
[2]: import sys
sys.path.append('/content/gdrive/My Drive/RL/.')
# sys.path.append('</dir/to/start/code/folder/>')
```

Your code should remain in the block marked by #####
YOUR CODE STARTS HERE # YOUR CODE ENDS HERE
Please don't edit anything outside the block.

```
[3]: %load_ext autoreload
%autoreload 2
```

```
[4]: import numpy as np
import random
import matplotlib.pyplot as plt
import gym
```

0.3 1. Incremental Implementation of Average

We've finished the incremental implementation of average for you. Please call the function `estimate` with `1/step` step size and fixed step size to compare the difference between this two on a simulated Bandit problem.

```
[5]: from RLalgs.utils import estimate
random.seed(6885)
numTimeStep = 10000
q_h = np.zeros(numTimeStep + 1) # Q Value estimate with 1/step step size
q_f = np.zeros(numTimeStep + 1) # Q value estimate with fixed step size
FixedStepSize = 0.5 #A large number to exaggerate the difference
for step in range(1, numTimeStep + 1):
    if step < numTimeStep / 2:
        r = random.gauss(mu = 1, sigma = 0.1)
    else:
        r = random.gauss(mu = 3, sigma = 0.1)

    #TIPS: Call function estimate defined in ./RLalgs/utils.py
    #####
    # YOUR CODE STARTS HERE
    q_h[step] = estimate(q_h[step-1],1/step,r)
    q_f[step] = estimate(q_f[step-1],FixedStepSize,r)
    # YOUR CODE ENDS HERE
    #####


q_h = q_h[1:]
q_f = q_f[1:]
```

RLalgs is a package containing Reinforcement Learning algorithms Epsilon-Greedy, Policy Iteration, Value Iteration, Q-Learning, and SARSA.

Plot the two Q value estimates. (Please include a title, labels on both axes, and legends)

```
[6]: #####
# YOUR CODE STARTS HERE
y1 = q_h
y2 = q_f
plt.plot(np.arange(numTimeStep), y1, label = 'q_h')
plt.plot(np.arange(numTimeStep), y2, label = 'q_f')
plt.xlabel('Time Step')
plt.ylabel('Q Values')
plt.legend()
plt.show()

# YOUR CODE ENDS HERE
#####
```



0.4 2. ϵ -Greedy for Exploration

In Reinforcement Learning, we are always faced with the dilemma of exploration and exploitation. ϵ -Greedy is a trade-off between them. You are gonna implement Greedy and ϵ -Greedy. We combine these two policies in one function by treating Greedy as ϵ -Greedy where $\epsilon = 0$. Edit the function `epsilon_greedy` in `./RLalgs/utils.py`.

```
[7]: from RLalgs.utils import epsilon_greedy
np.random.seed(6885) #Set the seed to cancel the randomness
q = np.random.normal(0, 1, size = 5)
#####
# YOUR CODE STARTS HERE
greedy_action = epsilon_greedy(q, 0) #Use epsilon = 0 for Greedy
e_greedy_action = epsilon_greedy(q, 0.1) #Use epsilon = 0.1
# YOUR CODE ENDS HERE
#####
print('Values:')
print(q)
print('Greedy Choice =', greedy_action)
print('Epsilon-Greedy Choice =', e_greedy_action)
```

Values:

[0.61264537 0.27923079 -0.84600857 0.05469574 -1.09990968]

Greedy Choice = 0

Epsilon-Greedy Choice = 0

You should get the following results. Values: [0.61264537 0.27923079 -0.84600857 0.05469574 -1.09990968] Greedy Choice = 0

0.5 3. Frozen Lake Environment

```
[8]: env = gym.make('FrozenLake-v1')
```

0.5.1 3.1 Derive Q value from V value

Edit function `action_evaluation` in `./RLalgs/utils.py`. TIPS: $q(s, a) = \sum_{s', r} p(s', r | s, a)(r + \gamma v(s'))$

```
[9]: from RLalgs.utils import action_evaluation
v = np.ones(16)
q = action_evaluation(env = env.env, gamma = 1, v = v)
print('Action values:')
print(q)
```

Action values:

```
[[1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.      1.      1.      ]
 [1.      1.33333333 1.33333333 1.33333333]
 [1.      1.      1.      1.      ]]
```

You should get Q values all equal to one except at State 14

Pseudo-codes of the following four algorithms can be found on Page 80, 83, 130, 131 of the Sutton's book.

0.5.2 3.2 Model-based RL algorithms

```
[10]: from RLalgs.utils import action_evaluation, action_selection, render
```

0.5.3 3.2.1 Policy Iteration

Edit the function `policy_iteration` and relevant functions in `./RLalgs/pi.py` to implement the Policy Iteration Algorithm.

```
[11]: from RLalgs.pi import policy_iteration
V, policy, numIterations = policy_iteration(env = env.env, gamma = 1,
↳max_iteration = 500, theta = 1e-7)
print('State values:')
print(V)
print('Number of iterations to converge =', numIterations)
```

State values:

```
[0.82352774 0.8235272 0.82352682 0.82352662 0.82352791 0.
0.52941063 0.          0.82352817 0.82352851 0.76470509 0.
0.          0.88235232 0.94117615 0.          ]
```

Number of iterations to converge = 7

You should get values close to: State values: [0.82352774 0.8235272 0.82352682 0.82352662 0.82352791 0. 0.52941063 0. 0.82352817 0.82352851 0.76470509 0.0. 0.88235232 0.94117615 0.]

```
[52]: #Uncomment and run the following to evaluate your result, comment them when you
↳generate the pdf
#Q = action_evaluation(env = env.env, gamma = 1, v = V)
#policy_estimate = action_selection(Q)
#render(env, policy_estimate)
```

0.5.4 3.2.2 Value Iteration

Edit the function value_iteration and relevant functions in ./RLalgs/vi.py to implement the Value Iteration Algorithm.

```
[12]: from RLalgs.vi import value_iteration
V, policy, numIterations = value_iteration(env = env.env, gamma = 1,
↳max_iteration = 500, theta = 1e-7)
print('State values:')
print(V)
print('Number of iterations to converge =', numIterations)
```

State values:

```
[0.82352937 0.82352936 0.82352935 0.82352935 0.82352938 0.
0.52941174 0.          0.82352938 0.82352939 0.76470586 0.
0.          0.88235293 0.94117646 0.          ]
```

Number of iterations to converge = 500

You should get values close to: State values: [0.82352773 0.82352718 0.8235268 0.8235266 0.8235279 0. 0.52941062 0. 0.82352816 0.8235285 0.76470509 0.0. 0.88235231 0.94117615 0.]

```
[ ]: #Uncomment and run the following to evaluate your result, comment them when you
↳generate the pdf
#Q = action_evaluation(env = env.env, gamma = 1, v = V)
#policy_estimate = action_selection(Q)
#render(env, policy_estimate)
```

0.5.5 3.3 Model free RL algorithms

0.5.6 3.3.1 Q-Learning

Edit the function QLearning in ./RLalgs/ql.py to implement the Q-Learning Algorithm.

```
[14]: from RLalgs.ql import QLearning
Q = QLearning(env = env.env, num_episodes = 1000, gamma = 1, lr = 0.1, e = 0.1)
print('Action values:')
print(Q)
```

Action values:

```
[[0.27495168 0.11447515 0.1440366 0.09456157]
 [0.04787668 0.11030049 0.00534849 0.01721119]
 [0.02033644 0.10776446 0.03926204 0.0197736 ]
 [0.05673026 0.01105529 0.00598546 0.01150785]
 [0.28902598 0.06238357 0.12446953 0.0851048 ]
 [0.         0.         0.         0.         ]
 [0.06445695 0.00662628 0.1025657 0.01121817]
 [0.         0.         0.         0.         ]
 [0.07742679 0.06008138 0.02624229 0.31471038]
 [0.02322962 0.35242292 0.11929348 0.05763229]
 [0.31739597 0.15489291 0.07188574 0.06213676]
 [0.         0.         0.         0.         ]
 [0.         0.         0.         0.         ]
 [0.1093972 0.31346873 0.45350831 0.1788256 ]
 [0.23223038 0.25733667 0.27648376 0.60552523]
 [0.         0.         0.         0.         ]]
```

Generally, you should get non-zero action values on non-terminal states.

```
[ ]: #Uncomment the following to evaluate your result, comment them when you
      ↪ generate the pdf
      #env = gym.make('FrozenLake-v1')
      #policy_estimate = action_selection(Q)
      #render(env, policy_estimate)
```

0.5.7 3.3.2 SARSA

Edit the function SARSA in ./RLalgs/sarsa.py to implement the SARSA Algorithm.

```
[17]: from RLalgs.sarsa import SARSA
Q = SARSA(env = env.env, num_episodes = 1000, gamma = 1, lr = 0.1, e = 0.1)
print('Action values:')
print(Q)
```

Action values:

```
[[0.01795687 0.02711272 0.05195552 0.00949569]
 [0.01105018 0.01693622 0.0062744 0.0502714 ]
 [0.05563381 0.04957945 0.07203931 0.03871152]]
```

```
[0.03895344 0.03664733 0.02514024 0.06863669]
[0.06302563 0.01107141 0.01637326 0.00458328]
[0.          0.          0.          0.          ]
[0.05938421 0.05928103 0.05955968 0.03506446]
[0.          0.          0.          0.          ]
[0.0147726  0.04607198 0.02512533 0.0943063  ]
[0.03678232 0.19270621 0.04088879 0.01029397]
[0.22570166 0.0121902  0.07933813 0.00896049]
[0.          0.          0.          0.          ]
[0.          0.          0.          0.          ]
[0.0091253  0.20923752 0.02927684 0.1083489  ]
[0.15061143 0.31659146 0.4868246  0.34165985]
[0.          0.          0.          0.          ]]
```

Generally, you should get non-zero action values on non-terminal states.

```
[ ]: # Uncomment the following to evaluate your result, comment them when you
      → generate the pdf
#env = gym.make('FrozenLake-v1')
#policy_estimate = action_selection(Q)
#render(env, policy_estimate)
```

0.5.8 3.4 Human

You can play this game if you are interested. See if you can get the frisbee either with or without the model.

```
[ ]: from RLalgs.utils import human_play
      # Uncomment and run the following to play the game, comment it when you generate
      → the pdf
#env = gym.make('FrozenLake-v1')
#human_play(env)
```

0.6 4. Exploration VS. Exploitation

Try to reproduce Figure 2.2 (the upper one is enough) of the Sutton's book based on the experiment described in [Chapter 2.3](#).

```
[18]: # Do the experiment and record average reward acquired in each time step
      #####
      # YOUR CODE STARTS HERE

      eps = [0.1,0.01,0] # different epsilons for eps-greedy algorithm to match those
      → in figure
      num_band = 1000 # number of bandits
      num_arms = 10 # number of arms
      rewards = [[],[],[]] # considering 3 separate epsilon values, we create 3 empty
      → lists to save the averages
```

```

actual = np.random.normal(0,1,(num_band,num_arms)) # actual rewards for
↳selecting an action

for iter in range(len(eps)):
    est = np.zeros((num_band,num_arms)) # intialize estimated rewards
    num_pulls = np.zeros((num_band,num_arms)) # initialize number of times an arm
↳is pulled

    for pull in range(1,num_band+1): # pull for as many bandits as we have,
↳higher the more it matches initial graph
        r_now = [] # all rewards in this pull

        for i in range(num_band):
            if np.random.random() < eps[iter]: # decide whether to explore or exploit
↳based on a random number
                arm = np.random.randint(num_arms) #
↳explore
            else :
                arm = np.argmax(est[i]) # exploit

            r_now.append(np.random.normal(actual[i][arm],1)) # append current pull
↳rewards
            num_pulls[i][arm] = num_pulls[i][arm]+1 # iterate the chose number pull
↳number for average calculation
            est[i][arm] = est[i][arm] + (np.random.normal(actual[i][arm],1) -
↳est[i][arm])/num_pulls[i][arm] # update the estimated reward for use in
↳future pulls

        avg_reward = np.mean(r_now)
        rewards[iter].append(avg_reward)
# YOUR CODE ENDS HERE
#####

```

You should get curves similar to that in the book.

```

[19]: # Plot the average reward
#####
# YOUR CODE STARTS HERE

x = range(1,1001)
plt.plot(x, rewards[0], 'b', label = 'eps = 0.1')
plt.plot(x, rewards[1], 'r', label = 'eps = 0.01')
plt.plot(x, rewards[2], 'g', label = 'eps = 0')
plt.xlabel('Steps')
plt.ylabel('Average Reward')
plt.legend()
plt.show()

```



```
# YOUR CODE ENDS HERE
#####
```

